

# Introduction to Spring A

**Core idea:** Spring AOP allows us to apply common behavior—security, auditing, and performance measurement—across methods without mixing that behavior into every business method.

## 1. A Simple Business Method

Consider a student-management application:

```
@Service
public class StudentService {

    private final StudentRepository studentRepository;

    public StudentService(StudentRepository studentRepository) {
        this.studentRepository = studentRepository;
    }

    public Student createStudent(Student student) {
        return studentRepository.save(student);
    }
}
```

The responsibility of `createStudent()` is easy to understand:

```
Receive student
  ↓
Save student
  ↓
Return saved student
```

This is the **business logic** of the method.

A business method answers a domain-related question, such

- How is a student created?
- How is an order placed?
- How is money transferred?
- How is inventory updated?
- How is a payment processed?

The method is currently focused on one responsibility, so its intent is immediately visible.

---

## 2. Real Applications Need More Than Business Logic

A production application may also need to:

- Log when a method begins
- Log whether it succeeds or fails
- Measure execution time
- Check whether the caller has permission
- Record an audit event
- Start, commit, or roll back a transaction
- Collect application metrics
- Retry temporary failures
- Cache a result
- Mask sensitive information

These requirements are important, but they are not part of the core process of creating a student.

If all of them are added directly to `createStudent()`, the method becomes difficult to read because most of its code no longer represents student-related business logic.

---

### 3. Business Logic vs Supporting Logic

A typical application contains two broad categories of response

#### Business Logic

Business logic represents the rules and operations of the application

Examples:

- Creating a student
- Placing an order
- Processing a payment
- Updating inventory
- Approving a loan

Business logic differs from one feature or domain to another.

#### Supporting or Infrastructure Logic

Supporting logic helps the application operate safely, reliably, securely, and observably.

Examples:

- Logging
- Security checks
- Transaction management
- Auditing
- Performance monitoring
- Caching

This logic is usually not specific to students, orders, payments, or inventory. The same supporting behaviour may be required across many unrelated features.

---

### 4. What Is a Concern?

A **concern** is a responsibility or area of interest within an application.

Examples:

- Student management is one concern.
- Payment processing is another concern.
- Logging is another concern.
- Security is another concern.
- Transaction management is another concern.

## Core Business Concerns

Core business concerns normally belong to a particular feature or module.

```
Student concern → StudentService
Order concern  → OrderService
Payment concern → PaymentService
```

These concerns are organized **vertically**. Each feature has its own classes and methods.

## Cross-Cutting Concerns

Some concerns are required across many business modules.

```
      Logging
      ↓
StudentService ← Security → OrderService
      ↓
      PaymentService
      ↓
      Auditing
```

These concerns cut horizontally across otherwise unrelated features. They are therefore called **cross-cutting concerns**.

A concern becomes cross-cutting when the same kind of behaviour must be applied in many different parts of the application.

## 5. Common Examples of Cross-Cutting

### Logging

The application may need to record:

- Which method was called
- When it started
- Whether it completed successfully
- Whether it failed
- Which arguments were supplied

The same logging requirement may apply to student, order, payment, inventory, and user services.

### Security

Suppose only administrators are allowed to delete students.

A permission check must run before the delete operation. Similar checks may also be needed before refunding payments, changing user roles, or approving requests.

### Performance Measurement

Suppose we want to identify slow methods.

A timer must run before the method, and the elapsed time must be calculated after the method finishes. Without a reusable mechanism, this timer would have to be added manually to every method that needs monitoring.

### Auditing

Certain business operations may need a permanent record, including:

- Who performed the operation
- What operation was performed
- When it happened
- Whether it succeeded or failed

## Transaction Management

Database operations may require a transaction to begin before the method and then commit or roll back after the method completes.

## 6. The Two Major Problems: Scattering and Repetition

Cross-cutting concerns create two major architectural problems that are often implemented directly inside business classes.

### 6.1 Scattering

**Scattering** means that the implementation of one concern is spread across many locations.

For example, logging code may appear in:

```
StudentService  
OrderService  
PaymentService  
InventoryService  
UserService  
ReportService
```

The logging concern is now scattered throughout the application.

Imagine an application containing 20 service classes with five methods each. Logging, auditing, or performance code may have to be added and maintained in approximately 100 different methods.

If the logging format or auditing policy changes, many classes may need to be modified.

### Problems Caused by Scattering

- The same concern is repeated in many places.
- Developers may implement it differently.
- A required method may be missed.
- A small policy change may require edits across the entire codebase.

- Maintenance becomes expensive.

## 6.2 Tangling

**Tangling** means that a single method contains code belonging to unrelated concerns.

```
public Payment processPayment(PaymentRequest request) {

    // Logging concern
    log.info("Processing payment");

    // Security concern
    securityService.checkPermission("PROCESS_PAYMENT");

    // Performance concern
    long start = System.currentTimeMillis();

    try {
        // Business concern
        Payment payment = gateway.charge(request);

        // Auditing concern
        auditService.record("PAYMENT_COMPLETED");

        return payment;
    } catch (Exception exception) {
        // Logging concern
        log.error("Payment failed", exception);

        // Auditing concern
        auditService.record("PAYMENT_FAILED");

        throw exception;
    }
}
```

```
    } finally {  
        // Performance concern  
        metrics.record(  
            System.currentTimeMillis() - s  
        );  
    }  
}
```

This one method is responsible for:

- Business logic
- Logging
- Security
- Auditing
- Performance measurement
- Exception-related behaviour

The problem is not that these responsibilities are unnecessary. The problem is that they are implemented together in the same method.

### Problems Caused by Tangling

1. The business intent becomes difficult to see.
2. Supporting behaviour may be applied inconsistently.
3. Changes become more expensive.
4. Developers can forget mandatory behaviour.
5. Testing becomes more complicated.
6. The method becomes harder to maintain.

Compare the two versions below.

### Tangled Version

```

public Student createStudent(Student student)

    log.info("Method started");
    securityService.checkPermission("CREATE_ST

    long start = System.currentTimeMillis();

    try {
        Student result = studentRepository.save(student);

        auditService.record("STUDENT_CREATED");

        return result;

    } catch (Exception exception) {
        log.error("Method failed", exception);
        throw exception;

    } finally {
        metrics.record(
            System.currentTimeMillis() - start
        );
    }
}

```

## Focused Version

```

public Student createStudent(Student student) {
    return studentRepository.save(student);
}

```

The second version communicates the business operation immediately. The challenge is to keep the method this clean while still applying the required supporting behaviour.

## 7. First Attempt: Utility Methods

A common first solution is to centralize repeated code inside

```
public class LoggingUtil {  
  
    public static void logStart(String methodName) {  
        System.out.println(methodName + " started");  
    }  
  
    public static void logEnd(String methodName) {  
        System.out.println(methodName + " completed");  
    }  
}
```

The service can then call the utility methods:

```
public Student createStudent(Student student) {  
  
    LoggingUtil.logStart("createStudent");  
  
    Student result = studentRepository.save(student);  
  
    LoggingUtil.logEnd("createStudent");  
  
    return result;  
}
```

This is better than duplicating the complete logging implementation, but it does not solve the fundamental problem.

### The Invocation Is Still Scattered

Every business method must still contain calls such as:

```
LoggingUtil.logStart(...);
```

```
LoggingUtil.logEnd(...);
```

The logging implementation is centralized, but the decision a **when to invoke it** remains scattered across business method

## Developers Must Remember to Call It

Nothing automatically guarantees that every required metho

A developer may:

- Forget to call `logStart()`
- Call `logStart()` but forget `logEnd()`
- Use inconsistent method names or messages
- Handle failure differently in different methods

## Success and Failure Paths Still Require Boilerplate

```
LoggingUtil.logStart("createStudent");

try {
    Student result = studentRepository.save(student);

    LoggingUtil.logSuccess("createStudent");

    return result;
} catch (Exception exception) {

    LoggingUtil.logFailure(
        "createStudent",
        exception
    );

    throw exception;
}
```

The utility avoids repeating the internal logging implementation. The business method still contains control-flow boilerplate.

Utility methods centralize **what the concern does**, but not **concern is applied**.

## 8. Second Attempt: Inheritance

Another possibility is to place common behaviour inside a base class.

```
public abstract class BaseService {

    protected void logStart(String methodName) {
        System.out.println(methodName + " started");
    }

    protected void logEnd(String methodName) {
        System.out.println(methodName + " completed");
    }
}
```

A service can extend this class:

```
@Service
public class StudentService extends BaseService {

    public Student createStudent(Student student) {

        logStart("createStudent");

        Student result = studentRepository.save(student);

        logEnd("createStudent");

        return result;
    }
}
```

```
}  
}
```

## Why Inheritance Does Not Solve the Problem

Inheritance shares the implementation of `logStart()` and `logE` method must still invoke them manually.

It also creates additional design problems:

- Java supports only single class inheritance.
- A service may already need to extend another class.
- Unrelated classes are forced into an artificial inheritance hierarchy.
- Logging, security, auditing, and other unrelated concerns may get mixed into one large base class.
- The parent-child relationship does not naturally represent these concerns.

Inheritance provides access to shared behaviour; it does not automatically execute that behaviour around selected methods.

## 9. Third Attempt: Manual Wrappers

A stronger solution is the **wrapper** or **decorator** approach.

Start with an interface:

```
public interface StudentOperations {  
  
    Student createStudent(Student student);  
}
```

The real service implements the interface:

```
public class StudentService implements StudentOperations {  
  
    private final StudentRepository studentRepository;
```

```

    public StudentService(StudentRepository studentRepository) {
        this.studentRepository = studentRepository;
    }

    @Override
    public Student createStudent(Student student) {
        return studentRepository.save(student);
    }
}

```

Now create a wrapper that implements the same interface:

```

public class LoggingStudentService
    implements StudentOperations {

    private final StudentOperations target;

    public LoggingStudentService(StudentOperations target) {
        this.target = target;
    }

    @Override
    public Student createStudent(Student student) {

        long start = System.currentTimeMillis();

        System.out.println("Method started");

        try {
            return target.createStudent(student);
        } finally {
            System.out.println(
                "Execution time: "
                + (System.currentTimeMillis() - start) + "s");
        }
    }
}

```

```

    tart)
        );
    }
}
}

```

The caller receives the wrapper instead of directly receiving

```

Controller
  ↓
LoggingStudentService
  ↓
StudentService

```

The original business class remains clean because the wrapper surrounds the real method invocation.

The wrapper performs three fundamental actions:

```

Run supporting behaviour before
  ↓
Call the real business object
  ↓
Run supporting behaviour after

```

This is close to the basic execution model used by Spring AOP.

AOP did not invent the idea of wrapping behaviour. It provides infrastructure for creating and applying such wrappers systematically.

## 10. Limitations of Manual Wrappers

Manual wrappers can work in a small application, but they become expensive at scale.

### A Wrapper Is Needed for Every Service Type

A wrapper written for `StudentOperations` cannot automatically wrap unrelated service interfaces such as `OrderOperations` or `PaymentOperations`.

Separate wrappers may be required for:

- Student services
- Order services
- Payment services
- Inventory services
- User services

## Multiple Concerns Require Multiple Layers

Suppose a service needs:

- Security
- Logging
- Performance measurement
- Auditing

The object chain may become:

```
Controller
  ↓
SecurityStudentService
  ↓
LoggingStudentService
  ↓
PerformanceStudentService
  ↓
AuditingStudentService
  ↓
StudentService
```

## Manual Object Construction Becomes Complicated

Someone must create all wrappers in the correct order and inject the final outer wrapper wherever the service is needed.

## The Same Boilerplate Appears Again

Each wrapper contains similar delegation and exception-handling logic. This concern has moved out of the business class, but wrapper class maintenance remains manual.

---

## 11. What a Better Solution Must Provide

A scalable solution should satisfy four requirements.

### Requirement 1: Keep Common Behaviour in One Place

Logging, security, auditing, or performance logic should have a dedicated module rather than being repeated in business classes.

### Requirement 2: Select Where the Behaviour Should Apply

We need a way to express rules such as:

```
Apply to all service methods
```

```
Apply only to methods annotated with @TrackExecutionTime
```

```
Apply to every method whose name starts with save
```

### Requirement 3: Control When the Behaviour Should Run

Supporting logic may need to run:

- Before the method
- After the method
- Only after successful completion
- Only after an exception is thrown
- Around the complete method execution

### Requirement 4: Avoid Modifying Every Business Method

We do not want every method to contain infrastructure calls:

```
public Student createStudent(Student student)
    startLogging();
    startTimer();
    checkSecurity();

    // Business logic
}
```

The application should apply cross-cutting behaviour without requiring the caller or business method to manage it manually.

The mechanism should remain largely transparent to the business code and its caller.

## 12. Introduction to Aspect-Oriented Programming

This is the problem that **Aspect-Oriented Programming**, or **AOP**, is designed to solve.

AOP provides a way to modularize cross-cutting concerns instead of mixing them into business classes.

The essential idea is:

```
Keep business logic in business classes
      +
Keep cross-cutting logic in a separate module
      +
Define where and when that logic should apply
      =
Combined behaviour at runtime
```

With AOP, the business method can remain focused:

```
public Student createStudent(Student student) {
    return studentRepository.save(student);
}
```

```
}
```

Logging, security, performance measurement, or auditing can be implemented through separate AOP configuration.

## 13. OOP and AOP Solve Different Dimensions

Object-Oriented Programming and Aspect-Oriented Programming are not competitors. They organize different dimensions of an application.

### OOP Organizes Business Capabilities Vertically

```
StudentService
├─ createStudent
├─ updateStudent
└─ deleteStudent

OrderService
├─ placeOrder
├─ cancelOrder
└─ getOrder

PaymentService
├─ makePayment
├─ refundPayment
└─ getPayment
```

Each class groups behaviour related to a particular business capability.

### AOP Organizes Behaviour Horizontally

```
Logging
↓
StudentService ← Performance → OrderService
↓
PaymentService
↓
Auditing
```

AOP captures behaviour that spans several business capabilities.

OOP organizes the core business capabilities.  
AOP organizes behaviour that cuts across those capabilities.

AOP complements OOP by modularizing concerns that do one business class.

## 14. The Three Fundamental Questions in AOP

AOP can be understood through three questions.

### Question 1: What Should Happen?

Examples:

```
Log method execution
Measure execution time
Check permission
Create an audit record
```

The code that performs this additional action is called **advice**.

### Question 2: Where Should It Happen?

Examples:

```
All methods in StudentService
All service-layer methods
Methods whose names begin with save
Methods annotated with @TrackExecutionTime
```

The rule that selects the relevant method executions is called a **pointcut**.

### Question 3: When Should It Happen?

Examples:

```
Before the method
After the method
```

```
After successful completion
After an exception
Around the complete method execution
```

The timing determines the type of advice.

## A Simple Mental Model

```
What should happen? → Advice
Where should it apply? → Pointcut
When should it run? → Advice type
```

## 15. What Is an Aspect?

An **aspect** is a module that contains a cross-cutting policy.

It normally brings together:

- The cross-cutting action that must run
- The rule that decides where it should run
- The timing of that action

For example:

```
PerformanceAspect
├─ Which methods should be measured?
└─ How should their execution time be measured?
```

Another example:

```
AuditAspect
├─ Which operations require auditing?
├─ What information should be recorded?
└─ Should both success and failure be recorded?
```

An aspect therefore represents a complete cross-cutting concern as a separate module.

## 16. How Spring AOP Applies an Aspect

Spring AOP normally creates a **proxy** around a Spring-managed

### Target Object

The **target** is the real object containing the business logic.

```
StudentService target object
```

### Proxy Object

When AOP is applied, the caller normally receives a proxy instead of directly receiving the target.

```
StudentService proxy
```

The proxy looks compatible with the target from the caller's perspective, but it can intercept method calls and apply additional behaviour.

### Runtime Flow

```
Controller
  ↓
AOP proxy
  ↓
Cross-cutting logic before the method
  ↓
Actual StudentService method
  ↓
Cross-cutting logic after the method
  ↓
Controller receives the result
```

Conceptually, the proxy behaves like this:

```
public Student createStudent(Student student) {

    logBefore();
```

```

try {
    Student result =
        target.createStudent(student);

    logSuccess(result);

    return result;

} catch (Exception exception) {

    logFailure(exception);

    throw exception;

}
}

```

The actual business method remains unchanged. Spring creates and manages the interception layer instead of requiring us to write a separate wrapper for every service.

## 17. Manual Wrapper vs Spring AOP Proxy

| Manual Wrapper                                      | Spring AOP                                                |
|-----------------------------------------------------|-----------------------------------------------------------|
| Written manually for a specific service type        | Proxy infrastructure is created and managed by Spring     |
| Delegation code must be repeated                    | Interception logic is reusable                            |
| Wrappers must be assembled manually                 | Spring integrates proxies with the IoC container          |
| Selection logic is hard-coded into wrapper creation | Pointcuts define where advice applies                     |
| Multiple concerns create nested wrapper classes     | Multiple advices can be applied through AOP configuration |
| Becomes difficult to scale                          | Designed for systematic cross-cutting behaviour           |

## 18. Important Spring AOP Terms at a Glance

| Term                         | Meaning                                                               |
|------------------------------|-----------------------------------------------------------------------|
| <b>Concern</b>               | A responsibility or area of interest in the application               |
| <b>Cross-cutting concern</b> | A concern needed across multiple unrelated concerns                   |
| <b>Aspect</b>                | A module that represents a cross-cutting concern                      |
| <b>Advice</b>                | The action that runs, such as logging or timing                       |
| <b>Pointcut</b>              | The rule that selects where advice should run                         |
| <b>Target</b>                | The real Spring bean containing the business logic                    |
| <b>Proxy</b>                 | The wrapper object created by Spring to intercept calls to the target |

## 19. Final Mental Model

Without AOP:

```
Business method
├─ Logging
├─ Security
├─ Performance measurement
├─ Auditing
└─ Actual business logic
```

With AOP:

```
Business class
└─ Actual business logic

Separate aspect
├─ Logging/security/auditing logic
├─ Rule defining where it applies
└─ Rule defining when it runs

Spring proxy
└─ Combines both behaviours at runtime
```

The complete idea can be summarized as:

```
Caller
↓
Spring AOP proxy
↓
Advice runs
↓
Target business method runs
↓
Remaining advice runs
↓
Result returns to caller
```

## 20. Key Takeaways

- Business methods should primarily express business intent.
- Logging, security, auditing, transactions, and monitoring are common cross-cutting concerns.
- Implementing these concerns directly creates scattering and tangling.
- Utility methods centralize implementation but still require manual invocation.
- Inheritance shares code but does not automatically apply behaviour.
- Manual wrappers demonstrate the right basic idea but do not scale well.
- AOP keeps cross-cutting logic in separate modules.
- Advice defines what happens.
- A pointcut defines where it happens.
- The advice type determines when it happens.
- An aspect combines the cross-cutting policy.
- Spring AOP commonly uses proxies to intercept calls to Spring-managed objects.
- OOP organizes business capabilities vertically, while AOP organizes shared behaviour horizontally.